


step 1: load data into jupyter

```
In [1]: 1 import pandas as pd
        2 import matplotlib.pyplot as plt
        3 import numpy as np
        4
        5 %matplotlib inline
```

```
In [2]: 1 df = pd.read_excel('pima-data.xlsx')
        2 df.head()
```

```
Out[2]:
```


	num_preg	glucose_conc	diastolic_bp	thickness	insulin	bmi	diab_pred	age	diabetes_orig
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1



```
In [3]: 1 df.tail()
```

```
Out[3]:
```

	num_preg	glucose_conc	diastolic_bp	thickness	insulin	bmi	diab_pred	age	diabetes_orig
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0



```
In [4]: 1 len(df)
```

```
Out[4]: 768
```

step 2: clean data

2.a - let us find if there are any null values

```
In [5]: 1 df.isnull().values.any()
```


```
Out[5]: False
```

2.b - let us find duplicate columns or co-related columns

```
In [6]: 1 # frist check how many columns are there
```

```
In [7]: 1 def plot_corr(df, size=12):  
2     corr = df.corr() #pandas data frame correlation function  
3     fig, ax = plt.subplots(figsize=(size,size))  
4     ax.matshow(corr) #color code the rectangles by correlation values  
5     plt.xticks(range(len(corr.columns)), corr.columns)  
6     plt.yticks(range(len(corr.columns)), corr.columns)
```

```
In [8]: 1 # Lets call above fuction  
2 plot_corr(df)
```



```
In [9]: 1 # we are going to remove thickness,diabetes_orig,has_diabetes
```

```
In [10]: 1 del df['thickness']  
2 del df['has_diabetes']  
3 del df['diabetes_orig']
```

```
In [11]: 1 df.head()
```

```
Out[11]:
```

	num_preg	glucose_conc	diastolic_bp	insulin	bmi	diab_pred	age	skin	diabetes
0	6	148	72	0	33.6	0.627	50	1.3790	True
1	1	85	66	0	26.6	0.351	31	1.1426	False
2	8	183	64	0	23.3	0.672	32	0.0000	True
3	1	89	66	94	28.1	0.167	21	0.9062	False
4	0	137	40	168	43.1	2.288	33	1.3790	True

2.c - lets convert text to numbers

machine learning algorithms will not understand text. so conver to numbers.

```
In [12]: 1 diabetes_map = {True:1, False:0}  
2 df['diabetes'] = df['diabetes'].map(diabetes_map)  
3
```

```
In [13]: 1 df.head()
```

```
Out[13]:
```

	num_preg	glucose_conc	diastolic_bp	insulin	bmi	diab_pred	age	skin	diabetes
0	6	148	72	0	33.6	0.627	50	1.3790	1
1	1	85	66	0	26.6	0.351	31	1.1426	0
2	8	183	64	0	23.3	0.672	32	0.0000	1
3	1	89	66	94	28.1	0.167	21	0.9062	0
4	0	137	40	168	43.1	2.288	33	1.3790	1

lets check proportion of diabetes vs non - diabetes data

we need to ensure that the proportions should be balanced(50-50) diabetes and non diabetes pr at least close enough to proceed

```
In [14]: 1 num_true = len(df.loc[df['diabetes'] == True]) #loc returns no of items with
2 num_false = len(df.loc[df['diabetes'] == False])
3 print('true = ', (num_true/ (num_true+num_false) )*100 ) #34% is decent figu
4 print('false = ', (num_false/ (num_true+num_false) )*100 ) #65% is decent fi
5
```

```
true = 34.89583333333333
false = 65.10416666666666
```

in case of data imbalancing we use SMOTE technique to increase lesser data samples

step 3 - train test split

let us split our data for training and testing the algorithm

```
In [15]: 1 from sklearn.model_selection import train_test_split
2 feature_col_names = ['num_preg','glucose_conc', 'diastolic_bp', 'skin','insu
3 'bmi','diab_pred','age']
4 predicted_class_names = ['diabetes']
5 x = df[feature_col_names].values #predictor feature columns (8xm)
6 y = df[predicted_class_names].values #predicted class [1=true, 0=false] colu
7 split_test_size = 0.30
8 x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=split_tes
9
```

```
In [16]: 1 print('# rows in dataframe {0}'.format(len(df)))
2 print('# rows missing glucose_conc : {0}'.format(len(df.loc[df['glucose_conc
3 == 0])))
4 print('# rows missing diastolic_bp : {0}'.format(len(df.loc[df['diastolic_bp
5 == 0])))
6 print('# rows missing thickness : {0}'.format(len(df.loc[df['skin']== 0
7 ])))
8 print('# rows missing insulin : {0}'.format(len(df.loc[df['insulin']== 0])))
9 print('# rows missing bmi : {0}'.format(len(df.loc[df['bmi']== 0])))
10 print('# rows missing diab_pred : {0}'.format(len(df.loc[df['diab_pred']== 0
11 ])))
12 print('# rows missing age : {0}'.format(len(df.loc[df['age']== 0])))
13
```

```
# rows in dataframe 768
# rows missing glucose_conc : 5
# rows missing diastolic_bp : 35
# rows missing thickness : 227
# rows missing insulin : 374
# rows missing bmi : 11
# rows missing diab_pred : 0
# rows missing age : 0
```

```
In [17]: 1 # Lets fill 0's with valid data (either mean or mode)
```

```
In [18]: 1 from sklearn.impute import SimpleImputer
2 fill_0 = SimpleImputer(missing_values = 0, strategy='mean')
3 x_train = fill_0.fit_transform(x_train)
4 x_test = fill_0.fit_transform(x_test)
```

```
In [19]: 1 x_train[0:10]
```

```
Out[19]: array([[1.00000000e+00, 9.50000000e+01, 6.00000000e+01, 7.09200000e-01,
5.80000000e+01, 2.39000000e+01, 2.60000000e-01, 2.20000000e+01],
[5.00000000e+00, 1.05000000e+02, 7.20000000e+01, 1.14260000e+00,
3.25000000e+02, 3.69000000e+01, 1.59000000e-01, 2.80000000e+01],
[4.34056399e+00, 1.35000000e+02, 6.80000000e+01, 1.65480000e+00,
2.50000000e+02, 4.23000000e+01, 3.65000000e-01, 2.40000000e+01],
[4.00000000e+00, 1.31000000e+02, 6.80000000e+01, 8.27400000e-01,
1.66000000e+02, 3.31000000e+01, 1.60000000e-01, 2.80000000e+01],
[1.00000000e+00, 1.03000000e+02, 3.00000000e+01, 1.49720000e+00,
8.30000000e+01, 4.33000000e+01, 1.83000000e-01, 3.30000000e+01],
[2.00000000e+00, 8.20000000e+01, 5.20000000e+01, 8.66800000e-01,
1.15000000e+02, 2.85000000e+01, 1.69900000e+00, 2.50000000e+01],
[3.00000000e+00, 1.28000000e+02, 7.80000000e+01, 1.12871227e+00,
1.55333333e+02, 2.11000000e+01, 2.68000000e-01, 5.50000000e+01],
[1.00000000e+00, 1.22000000e+02, 6.40000000e+01, 1.26080000e+00,
1.56000000e+02, 3.51000000e+01, 6.92000000e-01, 3.00000000e+01],
[4.34056399e+00, 1.38000000e+02, 7.22413127e+01, 1.12871227e+00,
1.55333333e+02, 3.63000000e+01, 9.33000000e-01, 2.50000000e+01],
[4.34056399e+00, 1.25000000e+02, 6.80000000e+01, 1.12871227e+00,
1.55333333e+02, 2.47000000e+01, 2.06000000e-01, 2.10000000e+01]])
```

step 4 - TRAIN THE MODEL

```
In [20]: 1 from sklearn.naive_bayes import GaussianNB
2 #create gaussian naive bayes model object and train it with data
3 nb_model = GaussianNB()
4 nb_model.fit(x_train, y_train.ravel())
```

```
Out[20]: GaussianNB()
```

```
In [21]: 1 # Lets test the algorithm's accuracy with training data it self.
```

```
In [22]: 1 nb_predict_train = nb_model.predict(x_train)
2 from sklearn import metrics
3 print('accuracy : {0:.4f}'.format(metrics.accuracy_score(y_train, nb_predict
4
```

```
accuracy : 0.7542
```

step 5 - TESTING MODEL

```
In [23]: 1 nb_predict_test = nb_model.predict(x_test)
2 from sklearn import metrics
3 print('accuracy : {0:.4f}'.format(metrics.accuracy_score(y_test, nb_predict_
4
```

accuracy : 0.7359

step 6 - ANALYZE THE MODEL ACCURACY

WITH THE HELP OF "CONFUSION MATRIX" WE CAN ANALYZE ALGORITHMS PERFORMANCE

```
In [24]: 1 print('confusion matrix')
2 print('{0}'.format(metrics.confusion_matrix(y_test, nb_predict_test)))
3 print('')
4 print('classification report')
5 print(metrics.classification_report(y_test, nb_predict_test))
```

confusion matrix

```
[[118  33]
 [ 28  52]]
```

classification report

	precision	recall	f1-score	support
0	0.81	0.78	0.79	151
1	0.61	0.65	0.63	80
accuracy			0.74	231
macro avg	0.71	0.72	0.71	231
weighted avg	0.74	0.74	0.74	231

final observation of navive bayes algo ==> accuracy 73%,recall 81%

type 2 error values should be less compared to type 1 error

```
In [25]: 1 # Let us try random forest algorithm
```

```
In [26]: 1 from sklearn.ensemble import RandomForestClassifier
2 rf_model = RandomForestClassifier(random_state=42)
3 rf_model.fit(x_train, y_train.ravel())
4 rf_predict_test = rf_model.predict(x_test)
5
6 #training metrics
7 print('accuracy : {0:.4f}'.format(metrics.accuracy_score(y_test, rf_predict_
8
```

accuracy : 0.7403

```
In [27]: 1 # Let us see confusion matrix and classification report
```

```
In [28]: 1 print(metrics.confusion_matrix(y_test, rf_predict_test))
2 print('')
3 print('classification report')
4 print(metrics.classification_report(y_test, rf_predict_test))
5
```

```
[[119  32]
 [ 28  52]]
```

```
classification report
              precision    recall  f1-score   support

     0       0.81      0.79      0.80       151
     1       0.62      0.65      0.63        80

 accuracy          0.74          231
 macro avg       0.71      0.72      0.72          231
 weighted avg    0.74      0.74      0.74          231
```

```
In [29]: 1 # Let us try with logistic regression algorithm
```

```
In [30]: 1 from sklearn.linear_model import LogisticRegression
2 lr_model = LogisticRegression(class_weight='balanced', C=0.2, random_state =
3 lr_model.fit(x_train, y_train.ravel())
4 lr_predict_test = lr_model.predict(x_test)
5
6 print('accuracy: {:.4f}'.format(metrics.accuracy_score(y_test, lr_predict_t
7 print('')
8 print(metrics.confusion_matrix(y_test, lr_predict_test))
9 print('')
10 print('classification report')
11 print(metrics.classification_report(y_test, lr_predict_test))
12
```

accuracy: 0.7143

```
[[111  40]
 [ 26  54]]
```

```
classification report
              precision    recall  f1-score   support

         0           0.81         0.74         0.77         151
         1           0.57         0.68         0.62          80

 accuracy                   0.71         231
 macro avg           0.69         0.71         0.70         231
weighted avg           0.73         0.71         0.72         231
```

C:\Users\91918\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

final conclusion - we are suggesting random forest lgorithm for this project as the accuracy is 75% and recall value is 82% which is higher than other two algorithms

```
In [ ]: 1
```